

# CS 640: Computer Networking

Ashutosh Shukla

Lecture 3  
Network Programming

---

---

---

---

---

---

---

## Topics

- Client-server model
- Sockets interface
- Socket primitives
- Example code for echoclient and echoserver
- Debugging With GDB
- Programming Assignment 1 (MNS)

---

---

---

---


---

---

---

## Client/sever model

- Client asks (*request*) – server provides (*response*)
- Typically: single server - multiple clients
- The server does not need to know *anything* about the client
  - even that it exists
- The client should always know *something* about the server
  - at least where it is located



*Note: clients and servers are processes running on hosts (can be the same or different hosts).*

---

---

---


---

---

---

---

## Sockets as means for inter-process communication (IPC)



The interface that the OS provides to its networking subsystem

---

---

---

---

---


---

---

---

## Internet Connections (TCP/IP)

- Address the machine on the network
  - By IP address
- Address the process
  - By the "port"-number
- The pair of IP-address + port – makes up a "socket-address"




---

---

---

---

---

---

---

---

## Clients

- Examples of client programs
  - Web browsers, ftp, telnet, ssh
- How does a client find the server?
  - The IP address in the server socket address identifies the host
  - The (well-known) port in the server socket address identifies the service, and thus implicitly identifies the server process that performs that service.
  - Examples of well known ports
    - Port 7: Echo server
    - Port 23: Telnet server
    - Port 25: Mail server
    - Port 80: Web server

---

---

---

---


---

---

---

---

## Using Ports to Identify Services



---

---

---

---

---

---

---

## Servers

- Servers are long-running processes (daemons).
  - Created at boot-time (typically) by the init process (process 1)
  - Run continuously until the machine is turned off.
- Each server waits for requests to arrive on a well-known port associated with a particular service.
  - Port 7: echo server
  - Port 23: telnet server
  - Port 25: mail server
  - Port 80: HTTP server
- Other applications should choose between 1024 and 65535

See /etc/services for a comprehensive list of the services available on a Linux machine.

---

---

---

---

---

---

---

## Sockets

- What is a socket?
  - To the kernel, a socket is an endpoint of communication.
  - To an application, a socket is a file descriptor that lets the application read/write from/to the network.
    - Remember: All Unix I/O devices, including networks, are modeled as files.
- Clients and servers communicate with each by reading from and writing to socket descriptors.
- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors.

---

---

---

---

---

---

---

## Socket Programming Cliches

- Network Byte Ordering
  - Network is big-endian, host may be big- or little-endian
  - Functions work on 16-bit (short) and 32-bit (long) values
  - htons() / htonl() : convert host byte order to network byte order
  - ntohs() / ntohl() : convert network byte order to host byte order
  - Use these to convert network addresses, ports, ...

```
struct sockaddr_in serveraddr;
/* fill in serveraddr with an address */
...
/* Connect takes (struct sockaddr *) as its second argument */
connect(clientfd, (struct sockaddr *) &serveraddr,
        sizeof(serveraddr));
...
```

- Structure Casts
  - You will see a lot of 'structure casts'

---

---

---

---

---

---

---

---

## Socket primitives

- **SOCKET**: int socket(int domain, int type, int protocol);

- domain := AF\_INET (IPv4 protocol)
- type := (SOCK\_DGRAM or SOCK\_STREAM)
- protocol := 0 (IPPROTO\_UDP or IPPROTO\_TCP)
- returned: socket descriptor (sockfd), -1 is an error

- **BIND**: int bind(int sockfd, struct sockaddr \*my\_addr, int addrlen);

- sockfd - socket descriptor (returned from socket())
- my\_addr: socket address, struct sockaddr\_in is used
- addrlen := sizeof(struct sockaddr)

```
struct sockaddr_in {
    unsigned short sin_family; /* address family (always AF_INET) */
    unsigned short sin_port; /* port num in network byte order */
    struct in_addr sin_addr; /* IP addr in network byte order */
    unsigned char sin_zero[8]; /* pad to sizeof(struct sockaddr) */
};
```

---

---

---

---

---

---

---

---

- **LISTEN**: int listen(int sockfd, int backlog);
  - backlog: how many connections we want to queue
- **ACCEPT**: int accept(int sockfd, void \*addr, int \*addrlen);
  - addr: here the socket-address of the caller will be written
  - returned: a new socket descriptor (for the temporal socket)
- **CONNECT**: int connect(int sockfd, struct sockaddr \*serv\_addr, int addrlen); //used by TCP client
  - parameters are same as for bind()
- **SEND**: int send(int sockfd, const void \*msg, int len, int flags);
  - msg: message you want to send
  - len: length of the message
  - flags := 0
  - returned: the number of bytes actually sent
- **RECEIVE**: int recv(int sockfd, void \*buf, int len, unsigned int flags);
  - buf: buffer to receive the message
  - len: length of the buffer ("don't give me more!")
  - flags := 0
  - returned: the number of bytes received

---

---

---

---

---

---

---

---

- **SEND** (DGRAM-style): `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
  - *msg*: message you want to send
  - *len*: length of the message
  - *flags* := 0
  - *to*: socket address of the remote process
  - *tolen*: = sizeof(struct sockaddr)
  - *returned*: the number of bytes actually sent
- **RECEIVE** (DGRAM-style): `int recvfrom(int sockfd, void *buf, int len, unsigned int flags, struct sockaddr *from, int *fromlen);`
  - *buf*: buffer to receive the message
  - *len*: length of the buffer ("don't give me more!")
  - *from*: socket address of the process that sent the data
  - *fromlen*: = sizeof(struct sockaddr)
  - *flags* := 0
  - *returned*: the number of bytes received
- **CLOSE**: `close (socketfd);`

---

---

---

---

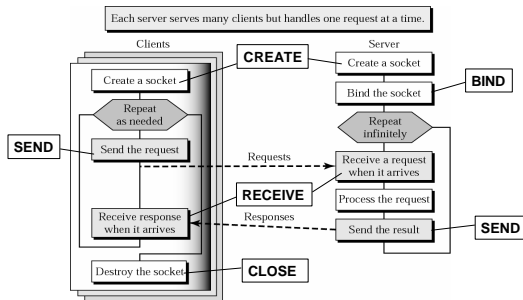
---

---

---

---

## Client+server: connectionless




---

---

---

---

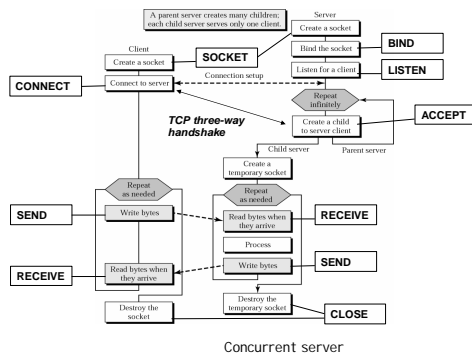
---

---

---

---

## Client+server: connection-oriented




---

---

---

---

---

---

---

---

### EchoClient.c – #include's

```
#include <stdio.h>    /* for printf() and fprintf() */
#include <sys/socket.h> /* for socket(), connect(),
                        sendto(), and recvfrom() */
#include <arpa/inet.h> /* for sockaddr_in and
                        inet_addr() */
#include <stdlib.h>    /* for atoi() and exit() */
#include <string.h>    /* for memset() */
#include <unistd.h>    /* for close() */

#define ECHOMAX 255   /* Longest string to echo */
```

---

---

---

---

---

---

---

### EchoClient.c -variable declarations

```
int main(int argc, char *argv[])
{
    int sock;          /* Socket descriptor */
    struct sockaddr_in echoServAddr; /* Echo server address */
    struct sockaddr_in fromAddr;    /* Source address of echo */
    unsigned short echoServPort=7; /* Echo server port */
    unsigned int fromSize;          /* address size for recvfrom() */
    char *servIP="172.24.23.4";     /* IP address of server */
    char *echoString="I hope this works"; /* String to send
    to echo server */
    char echoBuffer[ECHOMAX+1]; /* Buffer for receiving
    echoed string */
    int echoStringLen;          /* Length of string to echo */
    int respStringLen;          /* Length of received response */
```

---

---

---

---

---

---

---

### EchoClient.c - creating the socket and sending

```
/* Create a datagram/UDP socket */
sock = socket(AF_INET, SOCK_DGRAM, 0);

/* Construct the server address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out
structure */
echoServAddr.sin_family = AF_INET; /* Internet addr family */
echoServAddr.sin_addr.s_addr = htonl(servIP); /* Server IP
address */
echoServAddr.sin_port = htons(echoServPort); /* Server port */

/* Send the string to the server */
sendto(sock, echoString, echoStringLen, 0, (struct sockaddr *)
&echoServAddr, sizeof(echoServAddr));
/* Recv a response */
```

---

---

---

---

---

---

---

## EchoClient.c – receiving and printing

```
fromSize = sizeof(fromAddr);
recvfrom(sock, echoBuffer, ECHOMAX, 0, (struct sockaddr *)
    &fromAddr, &fromSize);
/* Error checks like packet is received from the same server*/

/* null-terminate the received data */
echoBuffer[echoStringLen] = '\0';
printf("Received: %s\n", echoBuffer); /* Print the echoed arg */
close(sock);
exit(0);
} /* end of main () */
```

---

---

---

---

---

---

---

## EchoServer.c

```
int main(int argc, char *argv[])
{
    int sock; /* Socket */
    struct sockaddr_in echoServAddr; /* Local address */
    struct sockaddr_in echoCIntAddr; /* Client address */
    unsigned int cliAddrLen; /* Length of incoming message */
    char echoBuffer[ECHOMAX]; /* Buffer for echo string */
    unsigned short echoServPort = 7; /* Server port */
    int recvMsgSize; /* Size of received message */
    /* Create socket for sending/receiving datagrams */
    sock = socket(AF_INET, SOCK_DGRAM, 0);
    /* Construct local address structure */
    memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out structure */
    echoServAddr.sin_family = AF_INET; /* Internet address family */
    echoServAddr.sin_addr.s_addr = htonl("172.24.23.4");
    echoServAddr.sin_port = htons(echoServPort); /* Local port */

    /* Bind to the local address */
    bind(sock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr);
```

---

---

---

---

---

---

---

```
for (;;) /* Run forever */
{
    cliAddrLen = sizeof(echoCIntAddr);

    /* Block until receive message from a client */
    recvMsgSize = recvfrom(sock, echoBuffer, ECHOMAX, 0,
        (struct sockaddr *) &echoCIntAddr, &cliAddrLen);

    printf("Handling client %s\n", inet_ntoa(echoCIntAddr.sin_addr));

    /* Send received datagram back to the client */
    sendto(sock, echoBuffer, recvMsgSize, 0,
        (struct sockaddr *) &echoCIntAddr, sizeof(echoCIntAddr);
}
} /* end of main () */
```

Error handling is must

---

---

---

---

---

---

---

## Socket Programming Help

- man is your friend
  - man accept
  - man sendto
  - Etc.
- The manual page will tell you:
  - What `#include<>` directives you need at the top of your source code
  - The type of each argument
  - The possible return values
  - The possible errors (in `errno`)

---

---

---

---

---

---

---

## Debugging with gdb

- Prepare program for debugging
  - Compile with `-g` (keep full symbol table)
  - Don't use compiler optimization (`-O0`, `-O2`, ...)
- Two main ways to run gdb
  - On program directly
    - `gdb progname`
    - Once gdb is executing we can execute the program with:
      - `run args`
  - On a core (post-mortem)
    - `gdb progname core`
    - Useful for examining program state at the point of crash
- Extensive in-program documentation exists
  - `help` (or `help <topic>` or `help <command>` )

---

---

---

---

---

---

---

## More information...

- Socket programming
  - W. Richard Stevens, UNIX Network Programming
  - Infinite number of online resources
  - <http://www.cs.rpi.edu/courses/sysprog/sockets/sock.html>
- GDB
  - Official GDB homepage:  
<http://www.gnu.org/software/gdb/gdb.html>
  - GDB primer: <http://www.cs.pitt.edu/~mosse/gdb-note.html>

---

---

---

---

---

---

---