

## CS 640: Introduction to Computer Networks

Aditya Akella

Lecture 24 -  
Peer-to-Peer

---

---

---

---

---

---

---

### The Road Ahead

- p2p file sharing techniques
  - Downloading: Whole-file vs. chunks
  - Searching
    - Centralized index (Napster, etc.)
    - Flooding (Gnutella, etc.)
    - Smarter flooding (Kazaa, ...)
    - Routing (Freenet, etc.)
- Uses of p2p - what works well, what doesn't?
  - servers vs. arbitrary nodes
  - Hard state (backups!) vs soft-state (caches)
- Challenges
  - Fairness, freeloading, security, ...

---

---

---

---

---

---

---

### P2p file-sharing

- Quickly grown in popularity
  - Dozens or hundreds of file sharing applications
  - 35 million American adults use P2P networks -- 29% of all Internet users in US!
  - Audio/Video transfer now dominates traffic on the Internet

---

---

---

---

---

---

---

## What's out there?

	Central	Flood	Super-node flood	Route
Whole File	Napster	Gnutella		Freenet
Chunk Based	BitTorrent		KaZaA (bytes, not chunks)	DHTs eDonkey 2000

---

---

---


---

---

---

---

## Publishing/Searching




---

---

---

---

---

---

---

## Searching

- **Needles vs. Haystacks**
  - Searching for top 40, or an obscure punk track from 1981 that nobody's heard of?
- **Search expressiveness**
  - Whole word? Regular expressions? File names? Attributes? Whole-text search?
    - (e.g., p2p gnutella or p2p google?)

---

---

---

---

---

---

---

## Framework

- Common Primitives:
  - **Join**: how to I begin participating?
  - **Publish**: how do I advertise my file?
  - **Search**: how to I find a file?
  - **Fetch**: how to I retrieve a file?

---

---

---

---

---

---

---

## Napster: Overview

- History
  - 1999: Sean Fanning launches Napster
  - Peaked at 1.5 million simultaneous users
  - Jul 2001: Napster shuts down
- Centralized Database:
  - **Join**: on startup, client contacts central server
  - **Publish**: reports list of files to central server
  - **Search**: query the server => return someone that stores the requested file
  - **Fetch**: get the file directly from peer

---

---

---


---

---

---

---

## Napster: Publish



---

---

---


---

---

---

---

## Napster: Search



---

---

---

---

---

---

---

## Napster: Discussion

- Pros:
  - Simple
  - Search scope is  $O(1)$
  - Controllable (pro or con?)
- Cons:
  - Server maintains lot of state
  - Server does all processing
  - Single point of failure

---

---

---

---

---

---

---

## Gnutella: Overview

- History:
  - In 2000, J. Frankel and T. Pepper from Nullsoft released Gnutella
  - Soon many other clients: Bearshare, Morpheus, LimeWire...
  - In 2001, many protocol enhancements including "ultrapeers"
- Query Flooding:
  - **Join**: on startup, client contacts a few other nodes; these become its "neighbors"
    - Ping-Pong protocol
  - **Publish**: no need
  - **Search**: ask neighbors, who ask their neighbors, and so on... when/if found, reply to sender.
    - TTL limits propagation
  - **Fetch**: get the file directly from peer

---

---

---


---

---

---

---

## Gnutella: Search




---

---

---

---

---

---

---

---

## Gnutella: Discussion

- **Pros:**
  - Fully de-centralized
  - Search cost distributed
  - Processing @ each node permits powerful search semantics
- **Cons:**
  - Search scope is  $O(N)$
  - Search time is  $O(???)$
  - Nodes leave often, network unstable
- **TTL-limited search works well for haystacks.**
  - For scalability, does NOT search every node. May have to re-issue query later

---

---

---

---

---

---

---

---

## KaZaA: Overview

- **Gnutella X Napster**
  - No dedicated server
  - But.. not all peers are equal!
- **"Smart" Query Flooding:**
  - **Join:** on startup, client contacts a "supernode" ... may at some point become one itself
  - **Publish:** send list of files to supernode
  - **Search:** send query to supernode, supernodes flood query amongst themselves.
  - **Fetch:** get the file directly from peer(s); can fetch simultaneously from multiple peers

---

---

---

---


---

---

---

---

## KaZaA: Network Design



---

---

---


---

---

---

---

## KaZaA: File Insert



---

---

---


---

---

---

---

## KaZaA: File Search



---

---

---

---

---

---

---

## KaZaA: Fetching

- More than one node may have requested file...
- How to tell?
  - Must be able to identify similar files with not necessarily same filename
  - Same filename not necessarily same file...
- Use Hash of file
  - KaZaA uses UUHash: fast, but not secure
  - Alternatives: MD5, SHA-1
- How to fetch?
  - Get bytes [0..1000] from A, [1001...2000] from B
  - Alternative: Erasure Codes

---

---

---

---

---

---

---

## Stability and Superpeers

- Why superpeers?
  - Query consolidation
    - Many connected nodes may have only a few files
    - Propagating a query to a sub-node would take more b/w than answering it yourself
  - Caching effect
    - Requires network stability
- Superpeer selection is time-based
  - How long you've been on is a good predictor of how long you'll be around.

---

---

---

---

---

---

---

## KaZaA: Discussion

- Pros:
  - Tries to take into account node heterogeneity:
    - Bandwidth
    - Host Computational Resources
    - Host Availability (?)
  - Rumored to take into account network locality
    - Why is this important?
- Cons:
  - Mechanisms easy to circumvent
    - Can freeloader easily
  - Still no real guarantees on search scope or search time
- Similar behavior to Gnutella, but better.

---

---

---

---

---

---

---

## BitTorrent: History

- Key Motivation:
  - Popularity exhibits temporal locality (Flash Crowds)
  - E.g., Slashdot effect, CNN on 9/11, new movie/game release
- Focused on Efficient *Fetching*, not *Searching*:
  - Distribute the *same* file to all *peers*
  - Single publisher, multiple downloaders
- Has some "real" publishers

---

---

---

---

---

---

---

## BitTorrent: Overview

- Swarming:
  - **Publish:** Run a tracker server.
  - **Search:** Out-of-band. E.g., use Google to find a tracker for the file you want.
  - **Join:** contact centralized "tracker" server, get a list of peers.
  - **Fetch:** Download chunks of the file from your peers. Upload chunks you have to them.
- Big differences from Napster:
  - Chunk based downloading
  - "few large files" focus
  - Anti-freeloading mechanisms

---

---

---


---

---

---

---

## BitTorrent: Publish/Join



---

---

---

---


---

---

---



## BitTorrent: Fetch



---

---

---

---

---

---

---

## BitTorrent: Sharing Strategy

- Employ "Tit-for-tat" sharing strategy
  - A is downloading from some other people
    - A will let the fastest N of those download from him
  - Be optimistic: occasionally let freeloaders download
    - Otherwise no one would ever start!
    - Also allows you to discover better peers to download from when they reciprocate

---

---

---

---

---

---

---

## BitTorrent: Discussion

- Pros:
  - Works reasonably well in practice
  - Gives peers incentive to share resources; avoids freeloaders
- Cons:
  - Pareto Efficiency relatively weak condition
  - Central tracker server needed to bootstrap swarm
  - (Tracker is a design choice, not a requirement, as you know from your projects. Could easily combine with other approaches.)

---

---

---

---

---

---

---

## P2P: Summary

- Many different styles; remember pros and cons of each
  - centralized, flooding, swarming, unstructured and structured routing
- Lessons learned:
  - Single points of failure are very bad
  - Flooding messages to everyone is bad
  - Underlying network topology is important
  - Not all nodes are equal
  - Need incentives to discourage freeloading
  - Privacy and security are important
  - Structure can provide theoretical bounds and guarantees

---

---

---

---

---

---

---